

# AA 274

# Principles of Robotic Autonomy

Combinatorial motion planning



**Stanford**  
University




# Today's lecture

- Aim
  - Introduction to motion planning & specifically, combinatorial motion planning
- Readings:
  - Bertsekas, *Dynamic Programming and Optimal Control, Vol I, 3<sup>rd</sup> ed.*, Section 2.3
  - LaValle, *Planning Algorithms*, Sections 6.1-6.3, 6.5



# Robot Motion Planning

- **Robot motion planning**: compute a sequence of actions that drives a robot from an initial condition to a terminal condition while avoiding obstacles, respecting kinematic/dynamical constraints, and possibly optimizing an objective function



# More examples of motion planning

- Game puzzles
- Steering autonomous vehicles
- Controlling humanoid robot
- Protein docking
- Surgery planning
- ...



# Some History



- Formally defined in the 1970s
- Development of exact, combinatorial solutions in the 1980s
- Development of sampling-based methods in the 1990s and deployment on real-time systems in the 2000s
- Current research: inclusion of differential constraints, planning under uncertainty, parallel implementation, feedback plans

## Popular Methods:

- Potential functions [Rimon, Koditschek, '92]
- Grid-based search ( $A^*$ ,  $D^*$ ) [Stentz, '94]
- Geometric algorithms (visibility graphs, cell decomposition) [LaValle, '06]
- Sampling-based [Kavraki et al, '96; LaValle, Kuffner, '06, ]



# Simplest Setup

- Assume 2D workspace:  $\mathcal{W} \subseteq \mathbb{R}^2$
- $\mathcal{O} \subset \mathcal{W}$  is the obstacle region with polygonal boundary
- Robot is a rigid polygon
- **Problem:** Given initial placement of robot, compute how to gradually move it into a desired goal placement so that it never touches the obstacle region



# Simplest Setup

**Key point:** motion planning problem described in the real-world, but it really lives in a another space -- the **configuration** (C-) space!



# Solution Approaches


## Three main approaches to motion planning:

- *Grid-based planning*: Discretize problem into grid and run a graph-search algorithm (Dijkstra,  $A^*$ , ...)
- *Combinatorial planning*: constructs structures in the  $C$ -space that discretely and completely capture all information needed to perform planning
- *Sampling-based planning*: uses collision detection algorithms to probe and incrementally search the  $C$ -space for a solution, rather than completely characterizing all of the  $C_{\text{free}}$  structure



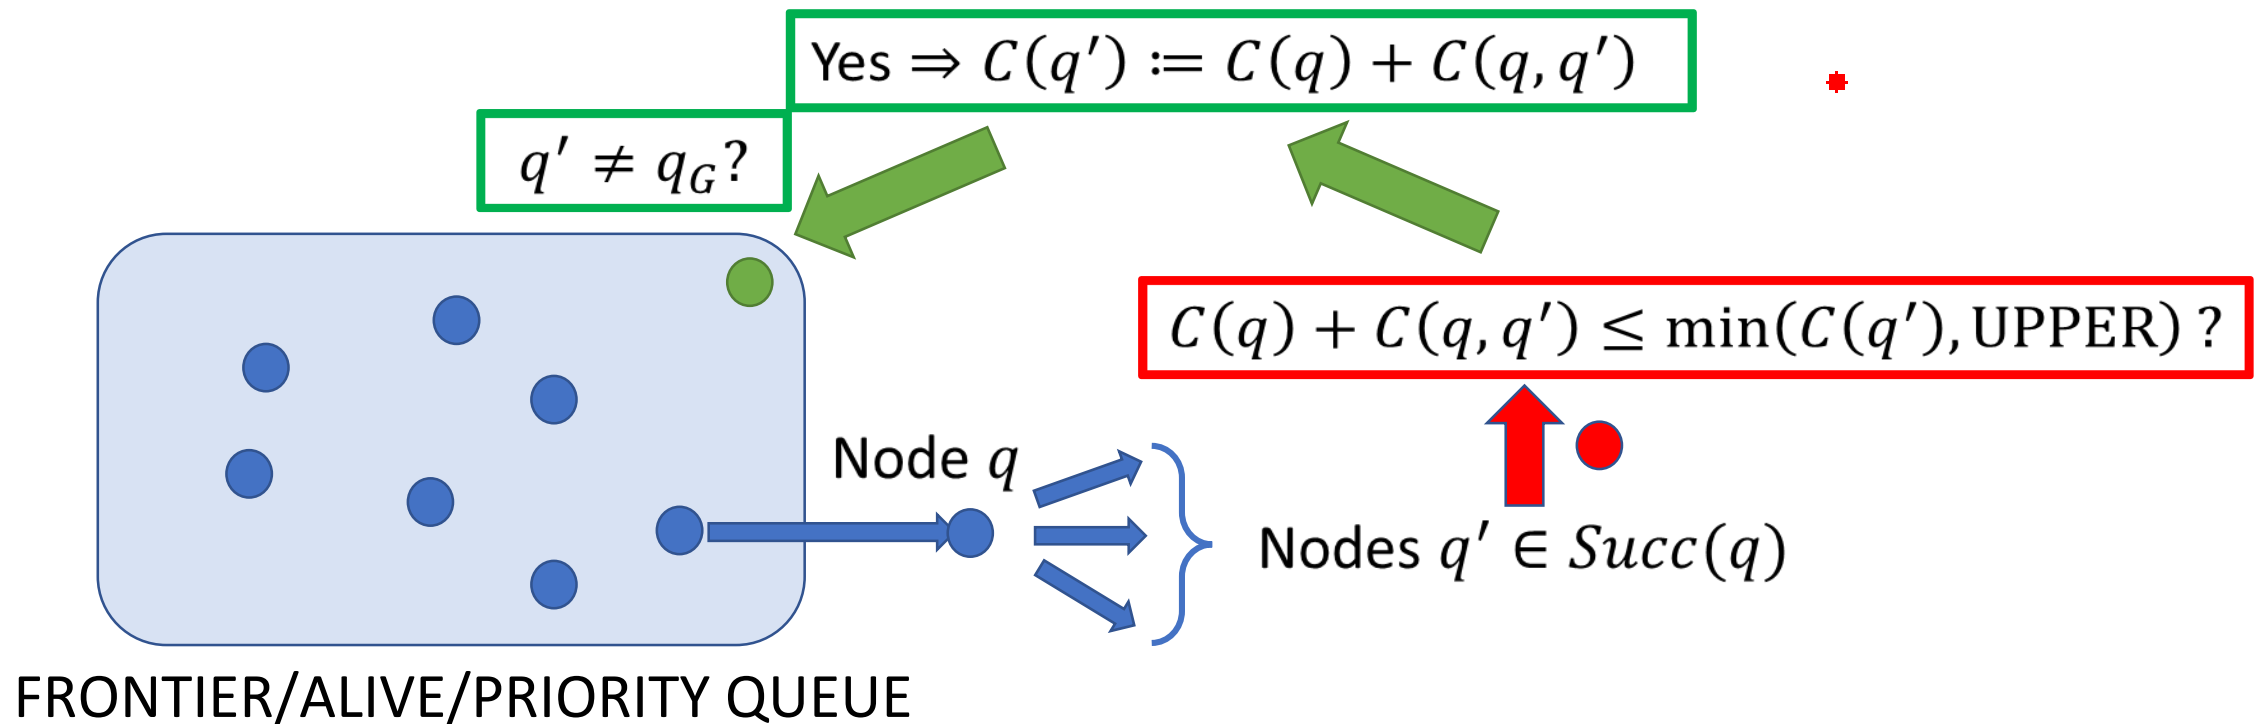
# Grid-based approaches

- Discretize the continuous world into a grid
  - Each grid cell is either free or forbidden
  - Robot moves between adjacent free cells
  - **Goal:** find a sequence of free cells from start to goal
- Mathematically, this corresponds to pathfinding in a discrete graph  $G = (V, E)$ 
  - Each vertex  $v \in V$  represents a free cell
  - Edges  $(v, u) \in E$  connect adjacent grid cells



# Graph Search Algorithms

- Having determined decomposition, how to find “best” path?
- **Label-Correcting Algorithms:**  $C(q)$ : *cost-of-arrival* from  $q_I$  to  $q$



\* Animation from Wikipedia

# Label Correcting Algorithm

**Step 1.** Remove a node  $q$  from frontier queue and for each child  $q'$  of  $q$ , execute step 2

**Step 2.** If  $C(q) + C(q, q') \leq \min(C(q'), \text{UPPER})$ , set  $C(q') := C(q) + C(q, q')$  and set  $q$  to be the parent of  $q'$ . In addition, if  $q' \neq q_G$ , place  $q'$  in the frontier queue if it is not already there, while if  $q' = q_G$ , set UPPER to the new value  $C(q) + C(q, q_G)$


**Step 3.** If the frontier queue is empty, terminate, else go to step 1.

**Initialization:** set the labels of all nodes to  $\infty$ , except for the label of the origin node, which is set to 0.

# GetNext() ?


**Depth-First-Search (DFS):** Maintain  $Q$  as a **stack** – Last in/first out

- Lower memory requirement (only need to store part of graph)



**Breadth-First-Search (BFS, Bellman-Ford):** Maintain  $Q$  as a **list** – First in/first first out

- Update cost for all edges up to current depth before proceeding to greater depth
- Can deal with negative edge (transition) costs




**Best-First (BF, Dijkstra):** Greedily select next  $q$ :  $q = \operatorname{argmin}_{q \in Q} C(q)$

- Node will enter the frontier queue at most *once*
- Requires costs to be non-negative

# Correctness & Improvements

## Theorem

If a feasible path exists from  $q_I$  to  $q_G$ , then algorithm terminates in finite time with  $C(q_G)$  equal to the optimal cost of traversal,  $C^*(q_G)$ .



# A\*: Improving Dijkstra

- Dijkstra orders by optimal “*cost-to-arrival*”
- Faster results if order by “*cost-to-arrival*”+ (approximate) “*cost-to-go*”
- That is, strengthen test

$$C(q) + C(q, q') \leq \text{UPPER}$$

to

$$C(q) + C(q, q') + h(q') \leq \text{UPPER}$$

where  $h(q)$  is heuristic for optimal cost-to-go (specifically, a positive *underestimate*)

- In this way, fewer nodes will be placed in the frontier queue
- This modification still guarantees that the algorithm will terminate with a shortest path

Dijkstra



A\*




# Graph-based approach: summary

- Pros:
  - Simple and easy to use
  - Fast (for some problems)
- Cons:
  - Resolution dependent
    - Not guaranteed to find solution if grid resolution is not small enough
  - Limited to simple robots (at most 3 DOFs)
    - Grid size is exponential in the number of DOFs

We shall move on to more “realistic” approaches...

# Back to (continuous) motion planning




- Our robot is a geometric entity operating in continuous space
- *Combinatorial techniques* for motion planning capture the structure of this continuous space
  - Particularly, the regions in which the robot is not in collision with obstacles
- Such approaches are typically complete
  - I.e., guaranteed to find a solution;
  - and sometimes even an optimal one







# Configuration Space

- A robot's **configuration** is a complete specification of the position of every point of the robot
- A robot's **degrees of freedom** (dof) is the smallest number  $d \geq 1$  of coordinates needed to represent its configuration
- The  $d$ -dimensional space  $\mathcal{C}$  containing all possible configurations of the robot is called the configuration space (C-space)
- For instance, a polygonal robot translating and rotating in the plane has  $d = 3$  DOFs and its configuration space is  $\mathbb{R}^2 \times \mathcal{S}$




# Configuration space

- The subset  $\mathcal{F} \subseteq \mathcal{C}$  of all collision free configurations is the **free space**



# Planning in $C$ -space


- Let  $R(q) \subset W$  denote set of points in the world occupied by robot when in configuration  $q$
- Robot in collision  $\Leftrightarrow R(q) \cap O \neq \emptyset$
- Accordingly, *free* space is defined as:  $C_{free} = \{q \in C | R(q) \cap O = \emptyset\}$
- Path planning problem in  $C$ -space: compute a **continuous** path:  
 $\tau: [0,1] \rightarrow C_{free}$   
with  $\tau(0) = q_I$  and  $\tau(1) = q_G$



# Combinatorial Planning

## Example:

- Point robot in the plane



**Key idea:** compute a roadmap, which is a graph in which each vertex is a configuration in  $\mathcal{C}_{\text{free}}$  and each edge is a path through  $\mathcal{C}_{\text{free}}$  that connects a pair of vertices

# Free-space roadmaps

Given a complete representation of the free space, we compute a roadmap that captures its connectivity.

A roadmap should preserve:

1. **Accessibility:** it is always possible to connect some  $q$  to the roadmap (e.g.,  $q_I \rightarrow s_1, q_G \rightarrow s_2$ )
  2. **Connectivity:** if there exists a path from  $q_I$  to  $q_G$ , there exists a path on the roadmap from  $s_1$  to  $s_2$
- **Key idea:** a roadmap provides a discrete representation of the continuous motion planning problem without losing any of the original connectivity information needed to solve it


# Cell Decomposition

Typical approach: **cell decomposition**. General requirements:

- Each cell should be easy to traverse (ideally convex)

- Decomposition should be easy to compute


- Adjacencies between cells should be straightforward to determine



# Computing a trapezoidal decomposition:


For every vertex (corner) of the forbidden space:

- Extend a vertical ray until it hits the first edge from top and bottom
- Implementation details:
  - Compute intersection points with all edges, and take the closest ones
  - More efficient approaches exists




# Other roadmaps

Maximum clearance (medial axis)



Minimum distance (visibility graph)




**Note:** No loss in optimality for a proper choice of discretization



# Visibility graph


- Allows to compute shortest collision free paths
- Connect by an edge every two vertices of the forbidden space that are *visible* from each other
  - The straight-line path between them is collision-free
- Given query points, connect them to the graph in a similar fashion



# Free-space computation


- The free space is not known in advance
- We need to compute this space given the ingredients
  - Robot representation, i.e., its shape (polygon, polyhedron, ...)
  - Representation of obstacles
- To achieve this we do the following:
  - Contract the robot into a point
  - In return, inflate (or stretch) obstacles by the shape of the robots
  - Also known as **Minkowski sum**


$$P \oplus Q := \{p + q | p \in P, q \in Q\}$$



# Higher Dimensions



- Extensions to higher dimensions is challenging  $\Rightarrow$  algebraic decomposition methods





# Additional resources on combinatorial planning

- Visualization of C-space for polygonal robot:  
<https://www.youtube.com/watch?v=SBFwgR4K1Gk>
- Algorithmic details for Minkowski sums and trapezoidal decomposition: de Berg et al., “Computational geometry: algorithms and applications”, 2008
- Implementation in C++:  
Computational Geometry Algorithms Library




# Combinatorial planning: summary

- These approaches are complete and even optimal in some cases
  - Do not discretize or approximate the problem
- Have theoretical guarantees on the running time
  - I.e., computational complexity is known
- Usually limited to small number of DOFs
  - Computationally intractable for many problems
- Problem specific: each algorithm applies to a specific type of robot/problem
- Difficult to implement: require special software to reason with geometric data structures (CGAL)

# Summary

- Search: Dijkstra (Uniform cost), A\* (cost-to-go heuristic)
- Combinatorial planning: discretize  $C$ -space as a graph, then search for shortest path
- Decomposition depends on cost function
- Popular method: cell decomposition – non intuitive in higher dimensions
- In general, motion planning is (PSPACE-)hard!

# Next time: sampling-based planning





# Backup: additional details for A\*

- Basic ingredients of A\*
  - $q_{init}$ : start vertex;  $q_{goal}$ : target vertex
  - OPEN: list of known vertices that have not been expanded yet
  - CLOSED: expanded vertices
  - $C(q, q')$ : real cost of edge from  $q$  to  $q'$
  - $C(q)$ : upper bound on “cost-to-come” from  $q_{init}$  to  $q$
  - $h(q)$ : lower bound on “cost-to-go” from  $q$  to  $q_{goal}$
  - $f(q) = C(q) + h(q)$ : upper bound on total cost

# A\* algorithm

1.  $C(q) = \infty, f(q) = \infty$  for all  $q$ ;
2.  $f(q_{init}) = h(q_{goal}); C(q_{init}) = 0$
3. OPEN = { $q_{init}$ }; CLOSED := { };
4. while (OPEN not empty)
5.      $q := q$  in OPEN that minimizes  $f(q)$
6.     if  $q == q_{goal}$  return path
7.     OPEN.remove( $q$ ); CLOSED.add( $q$ )
8.     for all  $q'$  in { $q' \mid (q, q') \text{ in } G, q' \text{ not in CLOSED}$ } // expansion
9.         OPEN.add( $q'$ )
10.        if  $C(q') \leq C(q) + C(q, q')$
11.            continue;
12.         $q'.parent = q; C(q') = C(q) + C(q, q')$
13.         $f(q') = C(q') + h(q')$
14. Return failure



# A\*: theory

- $h$  is **admissible** if for every  $q$ ,  $h(q)$  is at most the actual cost from  $q$  to  $q_{goal}$
- $h$  is **monotonic** if for every two adjacent vertices  $q, q'$ , it holds that  $h(q) \leq C(q, q') + h(q')$
- **Theorem:** If  $h$  is admissible and monotonic the algorithm returns the shortest path
- Common heuristic:
  - Euclidean distance in configuration space